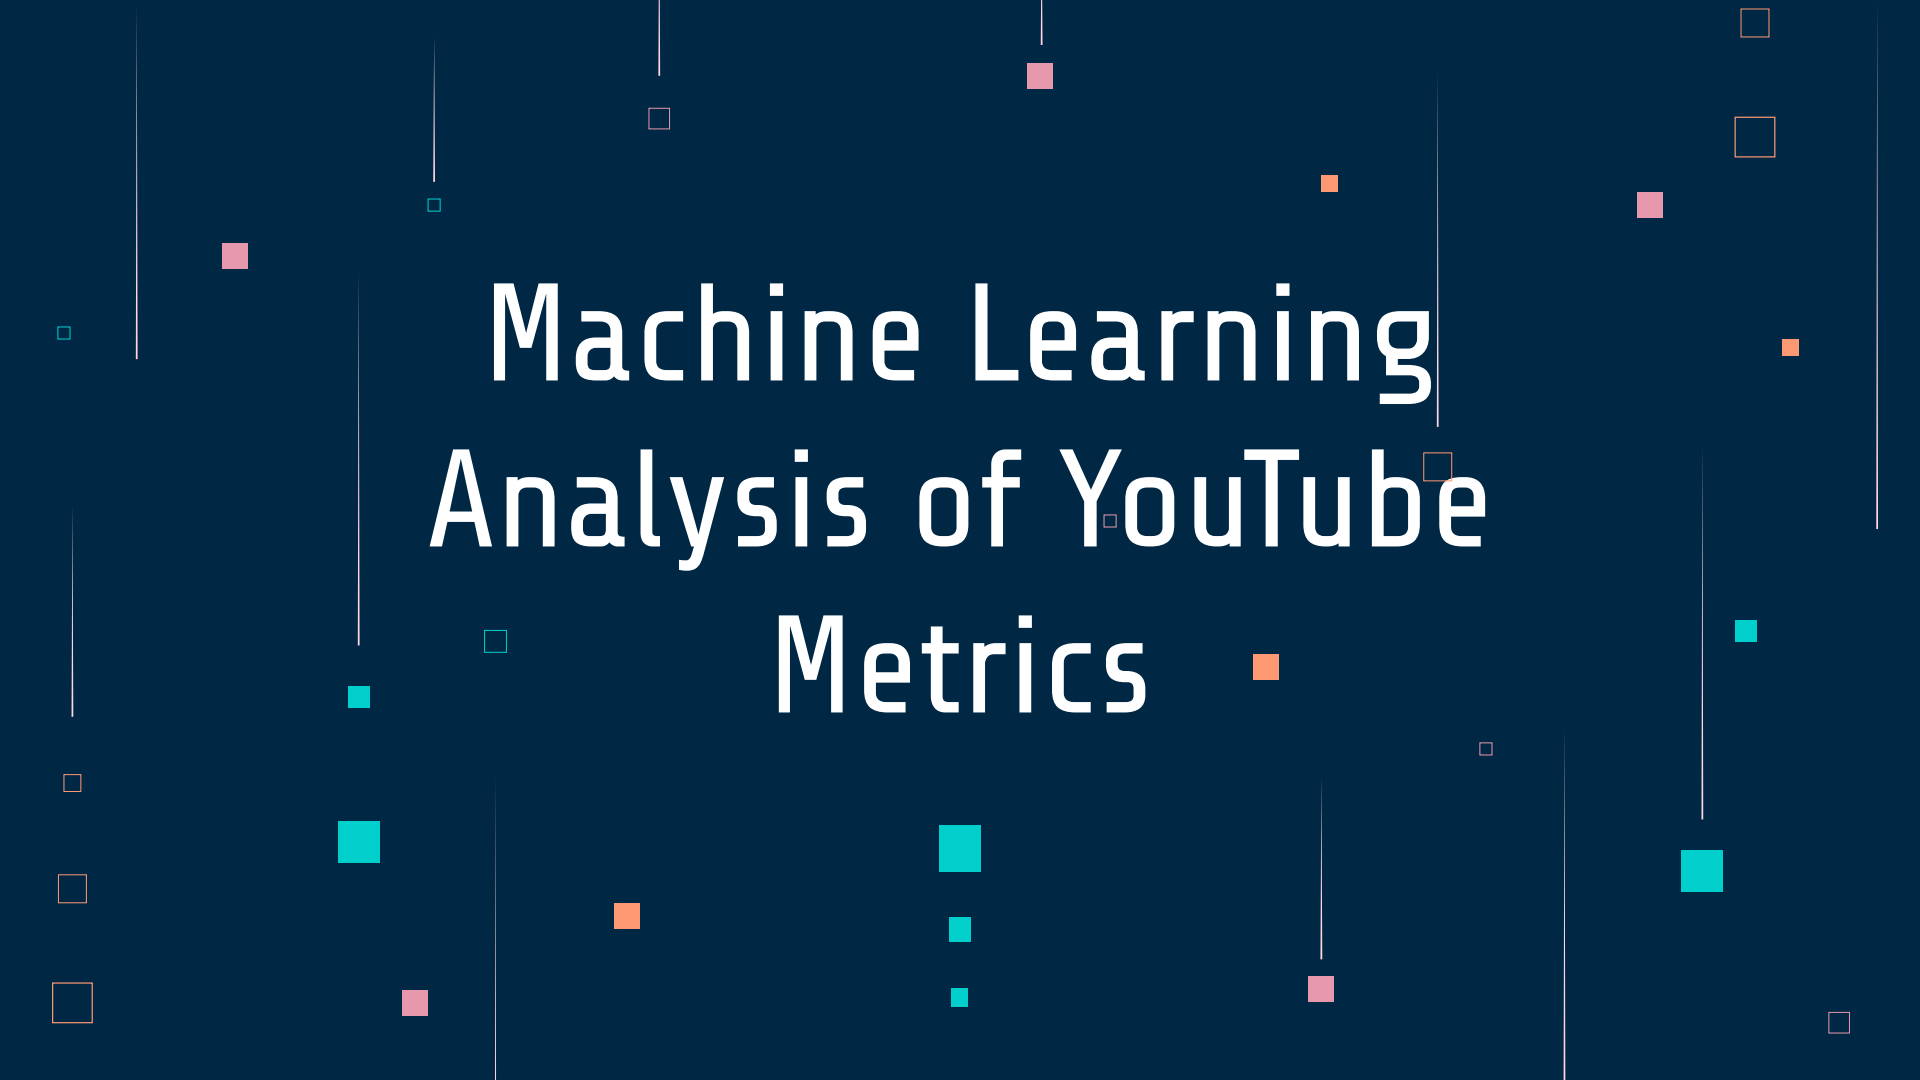


# Machine Learning Analysis of YouTube Metrics

The background of the slide is a dark blue gradient. It is decorated with an abstract pattern of small squares in various colors (pink, orange, teal, and light blue) and thin white vertical lines of varying lengths, scattered across the entire frame.

# OUR CONSULTANTS

Aaron Schneberger



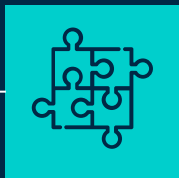
Ashok Goyal



Navyasri Pusuluri

Roli Singh


# TABLE OF CONTENTS



01

## PROBLEM & SOLUTION

We wanted to cluster YouTube videos and try to predict the revenue they generate



02

## OUR PROCESS

We used clustering algorithms and deep learning



03

## TARGET

We successfully clustered the videos and achieved a low mean percent error

# OBJECTIVES:

- To use unsupervised learning model “K-Means” to predict **clusters** of YouTube videos.
- To use supervised learning model “Keras” to predict the **revenue generated** by YouTube videos.
- To use supervised learning model “Keras” to predict the number of **views** of YouTube videos
- To **tune** the supervised learning model to find the best hyperparameters.

# Introducing the Data

The source data file from Kaggle includes metrics such as:

- Comments added: The number of comments on the video
- Views: The number of views the video has
- Shares: The number of times the video was shared
- Likes: The number of likes the video has
- Dislikes: The number of dislikes the video has
- RPM (USD): The revenue per thousand views the video has
- CPM (USD): The cost per thousand views the video has
- Watch time (hours): The total number of hours the video has been watched.

# Data Understanding/Pre-Processing

Data Extraction and Cleaning



Heat map

Scaling

PCA


K-Means Clustering


Agglomerative Clustering



# K Means Clustering/ Agglomerative Clustering

- Removed columns from data
- Standardized
- Made a heat map of features
- Used PCA to reduce dimensionality to 4 with a total explained variance of 0.9677





	PC1	PC2	PC3	PC4
<b>Comments added</b>	0.053539	0.615340	-0.203983	0.752476
<b>Shares</b>	0.083922	0.365346	0.307977	-0.272334
<b>Dislikes</b>	0.109409	0.203606	0.018685	-0.199867
<b>Likes</b>	0.078747	0.378704	0.315887	-0.256393
<b>RPM (USD)</b>	0.709812	0.074660	-0.623908	-0.280986
<b>CPM (USD)</b>	0.669248	-0.341088	0.541028	0.371087
<b>Views</b>	0.081488	0.300242	0.222544	-0.171631
<b>Watch time (hours)</b>	0.116416	0.301573	0.178632	-0.088401





# DEEP LEARNING REGRESSION

We changed the revenue per 1000 views to revenue then used it as the target vector:

```
In [9]: y = df["RPM (USD)"]*df["Views"]/1000
```

```
In [11]: X = df.drop(["RPM (USD)", "Views"], axis = 1)
```

# DEEP LEARNING REGRESSION

We used the other columns (other than views) as the features:

```
In [11]: X = df.drop(["RPM (USD)", "Views"], axis = 1)
```

```
In [12]: X.rename(columns = {"CPM (USD)": "Cost"}, inplace = True)
```

```
In [13]: X
```

Out[13]:

	Comments added	Shares	Dislikes	Likes	Cost	Watch time (hours)
1	907	9583	942	46903	16089.429765	65850.7042
2	412	4	4	130	14.339369	200.2966
3	402	152	15	881	249.688250	3687.3387
4	375	367	22	2622	393.686852	2148.3110
5	329	118	15	590	99.710325	1034.3945
...	...	...	...	...	...	...
218	4	5	0	30	46.287850	9.6188
219	3	5	1	48	15.874896	56.5930
220	3	0	0	44	8.546608	19.2752
221	3	1	0	35	9.077390	22.5450
222	2	5	0	38	13.353364	57.6363

# FIRST ATTEMPT AT REGRESSION

```
In [89]: # Define the model with six hidden layers and an output layer with one unit
nn = Sequential()

# First hidden layer
nn.add(Dense(units=8, activation="relu", input_dim=6))

for i in range(5):
    nn.add(Dense(units=8, activation="relu"))


# Output layer
nn.add(Dense(units=1, activation='relu'))

# Check the structure of the model
nn.summary()
```

## RESULTS OF FIRST ATTEMPT


- Test data had mean squared error of 7036.71
- Root mean error was 83.89
- The mean of the revenue column is \$132.21, and the max is \$7963.86

# DISTRIBUTION OF REVENUE



Root mean error  
of 83.89 quite  
high for this  
distribution

# LOSS VERSUS EPOCHS



We had some  
overfitting in our  
first attempt

# LOSS FUNCTION

- Decided that absolute percentage error was a better loss function

```
In [92]: # Evaluate the model using the test data
model_loss, model_percent_error = nn.evaluate(X_test_scaled,y_test,verbose=2)
print(f"Loss: {model_loss}, Mean Absolute Percent Error: {model_percent_error}")

2/2 - 0s - loss: 7036.7100 - mean_absolute_percentage_error: 424.5189 - 107ms/epoch - 54ms/step
Loss: 7036.7099609375, Mean Absolute Percent Error: 424.5188903808594
```



# KERAS TUNER

- Used keras tuner to find better hyperparameters

```
In [99]: # Create a method that creates a new Sequential model with hyperparameter options
def create_model_percent_error(hp):
    nn_model = Sequential()

    # Allow kerastuner to decide which activation function to use in hidden layers
    activation = hp.Choice('activation', ['relu', 'tanh'])

    # Allow kerastuner to decide number of neurons in first layer
    nn_model.add(Dense(units=hp.Int('first_units',
                                   min_value=1,
                                   max_value=15,
                                   step=5), activation=activation, input_dim=6))

    # Allow kerastuner to decide number of hidden layers and neurons in hidden layers
    for i in range(hp.Int('num_layers', 1, 7)):
        nn_model.add(Dense(units=hp.Int('units_' + str(i),
                                         min_value=1,
                                         max_value=15,
                                         step=5),
                           activation=activation))

    nn_model.add(Dense(units=1, activation="relu"))

    # Compile the model
    nn_model.compile(loss="mean_absolute_percentage_error", optimizer='adam', metrics = ['mean_absolute_percentage_error'])

    return nn_model
```

# RESULTS OF TUNING

Here are hyperparameters for best model

```
In [107]: # Find top model hyperparameters and print the values for mean absolute percentage error loss function
top_hyper_percent_loss = tuner.get_best_hyperparameters(1)
for param in top_hyper_percent_loss:
    print(param.values)

{'activation': 'relu', 'first_units': 6, 'num_layers': 4, 'units_0': 11, 'units_1': 11, 'units_2': 1, 'units_3': 11,
'units_4': 11, 'units_5': 6, 'units_6': 1, 'tuner/epochs': 200, 'tuner/initial_epoch': 67, 'tuner/bracket': 2, 'tuner/round': 2, 'tuner/trial_id': '0221'}
```

# RESULTS OF TUNING

Best model had mean absolute percentage error of 37.86 during validation

```
# Find the best model for mean absolute percentage error
top_model_mean_absolute_percentage_error = tuner.get_best_models(1)[0]


model_loss = top_model_mean_absolute_percentage_error.evaluate(X_test_scaled,y_test,verbose=1)
print(f"Mean Absolute Percentage Error: {model_loss}")

2/2 [=====] - 0s 4ms/step - loss: 37.8564 - mean_absolute_percentage_error: 37.8564
Mean Absolute Percentage Error: [37.856414794921875, 37.856414794921875]
```

But we felt there was overfitting


# RESULTS OF TUNING

We ran the  
training again and  
saw the loss graph  
to the right



# RESULTS OF TUNING

After using 175 for our number of epochs, we got a better loss curve:



# RESULTS OF TUNING

When evaluating the model, we got a mean error of  
32.74%

```
In [38]: # Evaluate the model using the test data
model_loss = nn_final.evaluate(X_test_scaled,y_test,verbose=1)
print(f"Mean Absolute Percent Error: {model_loss}")
```

```
2/2 [=====] - 0s 3ms/step - loss: 32.7416 - mean_absolute_percentage_error: 32.7416
Mean Absolute Percent Error: [32.741600036621094, 32.741600036621094]
```

# SUMMARY

- Clustered videos with k-means and agglomerative
- Created deep learning networks to predict revenue
- Achieved a mean percent error of 32.74%



# THANKS



CREDITS: This presentation template was created by [Slidesgo](#), including icons by [Flaticon](#), and infographics & images by [Freepik](#)